

Load the dataset safely, inspect structure, handle missing values.

```
import pandas as pd
import numpy as np

# Load dataset
data = pd.read_csv("dataset/healthcare-dataset-stroke-data.csv")

# Display first 5 rows
print(data.head())

# Dataset info
print(data.info())

# Check missing values
print(data.isnull().sum())

      id  gender  age  hypertension  heart_disease  ever_married \
0    9046    Male  67.0          0              1        Yes
1   51676  Female  61.0          0              0        Yes
2   31112    Male  80.0          0              1        Yes
3   60182  Female  49.0          0              0        Yes
4   1665  Female  79.0          1              0        Yes

      work_type Residence_type  avg_glucose_level    bmi
smoking_status \
0           Private            Urban       228.69  36.6  formerly
smoked
1  Self-employed            Rural       202.21    NaN    never
smoked
2           Private            Rural       105.92  32.5    never
smoked
3           Private            Urban       171.23  34.4
smokes
4  Self-employed            Rural       174.12  24.0    never
smoked

      stroke
0      1
1      1
2      1
3      1
4      1
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   id               5110 non-null   int64  
 1   gender            5110 non-null   object 
 2   work_type         5110 non-null   object 
 3  Residence_type    5110 non-null   object 
 4   avg_glucose_level 5110 non-null   float64
 5   bmi               5110 non-null   float64
 6   ever_married     5110 non-null   object 
 7   stroke            5110 non-null   int64  
 8   smoking_status    5110 non-null   object 
 9   hypertension      5110 non-null   int64  
 10  heart_disease    5110 non-null   int64  
 11  age               5110 non-null   float64
```

```

2   age          5110 non-null    float64
3   hypertension  5110 non-null    int64
4   heart_disease 5110 non-null    int64
5   ever_married   5110 non-null    object
6   work_type      5110 non-null    object
7   Residence_type 5110 non-null    object
8   avg_glucose_level 5110 non-null    float64
9   bmi            4909 non-null    float64
10  smoking_status 5110 non-null    object
11  stroke         5110 non-null    int64
dtypes: float64(3), int64(4), object(5)
memory usage: 479.2+ KB
None
id              0
gender          0
age              0
hypertension     0
heart_disease   0
ever_married    0
work_type        0
Residence_type  0
avg_glucose_level 0
bmi             201
smoking_status  0
stroke          0
dtype: int64

```

Understand whether the dataset is imbalanced.

```

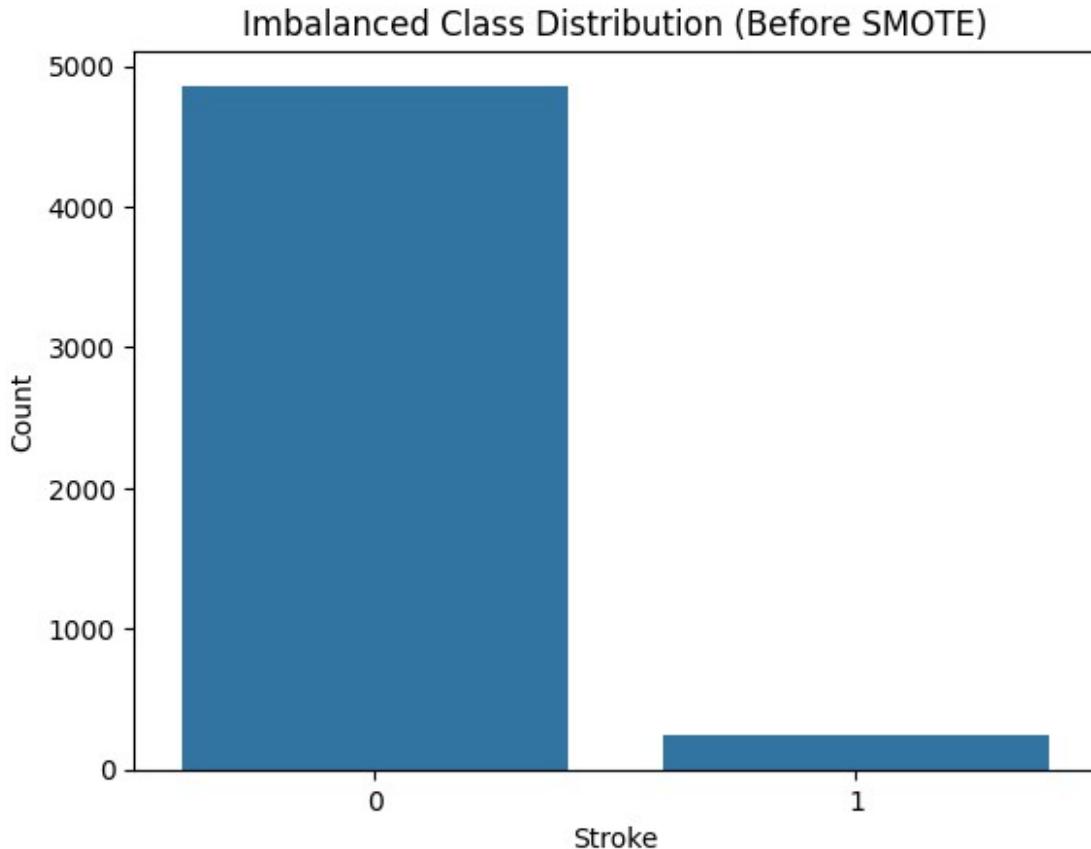
import matplotlib.pyplot as plt
import seaborn as sns

# Class distribution
class_counts = data['stroke'].value_counts()
print(class_counts)

# Visualization
plt.figure()
sns.countplot(x='stroke', data=data)
plt.title("Imbalanced Class Distribution (Before SMOTE)")
plt.xlabel("Stroke")
plt.ylabel("Count")
plt.show()

stroke
0    4861
1    249
Name: count, dtype: int64

```



Balance the dataset using Synthetic Minority Over-sampling Technique (SMOTE)

```

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from imblearn.over_sampling import SMOTE

# Drop unnecessary ID column
data = data.drop(columns=['id'])

# Handle missing BMI values
data['bmi'].fillna(data['bmi'].mean(), inplace=True)

# Encode categorical variables
label_cols = ['gender', 'ever_married', 'work_type', 'Residence_type',
'smoking_status']
le = LabelEncoder()

for col in label_cols:
    data[col] = le.fit_transform(data[col])

# Split features and target
X = data.drop('stroke', axis=1)
y = data['stroke']

```

```
# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# Apply SMOTE ONLY on training data
smote = SMOTE(random_state=42)
X_train_smote, y_train_smote = smote.fit_resample(X_train, y_train)

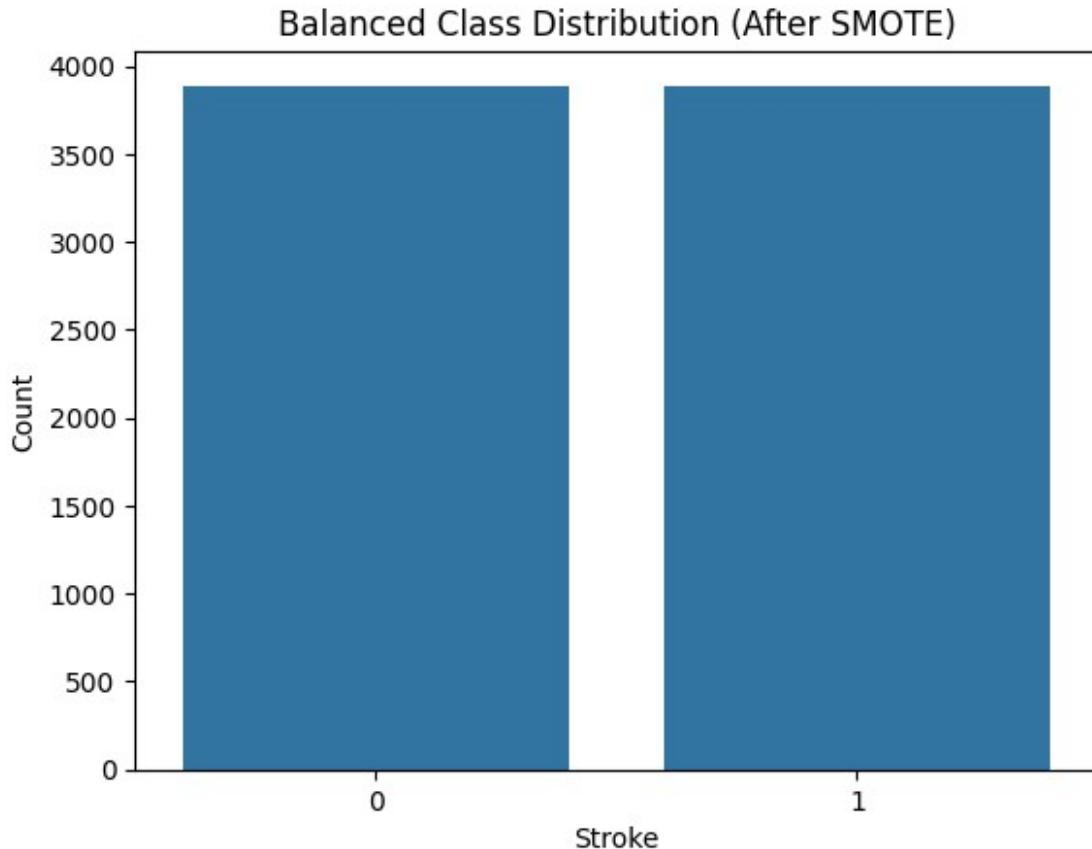
C:\Users\siddh\AppData\Local\Temp\ipykernel_24164\3244797451.py:9:
FutureWarning: A value is trying to be set on a copy of a DataFrame or
Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never
work because the intermediate object on which we are setting values
always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try
using 'df.method({col: value}, inplace=True)' or df[col] =
df[col].method(value) instead, to perform the operation inplace on the
original object.

data['bmi'].fillna(data['bmi'].mean(), inplace=True)
```

Visualization After SMOTE

```
plt.figure()
sns.countplot(x=y_train_smote)
plt.title("Balanced Class Distribution (After SMOTE)")
plt.xlabel("Stroke")
plt.ylabel("Count")
plt.show()
```



Classification Model

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix

# Train model
model = LogisticRegression(max_iter=1000)
model.fit(X_train_smote, y_train_smote)

# Predictions
y_pred = model.predict(X_test)
```

Evaluation Metrics

```
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

Confusion Matrix:
[[766 206]
 [ 14  36]]
```

Classification Report:					
	precision	recall	f1-score	support	
0	0.98	0.79	0.87	972	
1	0.15	0.72	0.25	50	
accuracy			0.78	1022	
macro avg	0.57	0.75	0.56	1022	
weighted avg	0.94	0.78	0.84	1022	